GA DSI 26 Project 3: Wine and Beer

By: Lim Zhi Yong

Task:

- Understand consumer patterns
- Identify if the consumer wants winemaking or homebrewing info
- Train model with subreddit posts

Data Description

- 1,000 posts from each subreddit
- Cleaned punctuation, stopwords, delimiters
- Considered both unigrams and bigrams

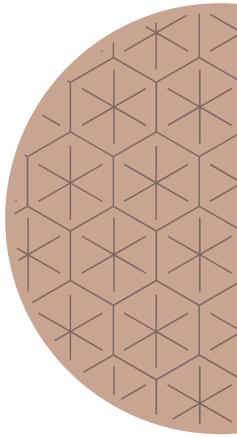

Model Workflow

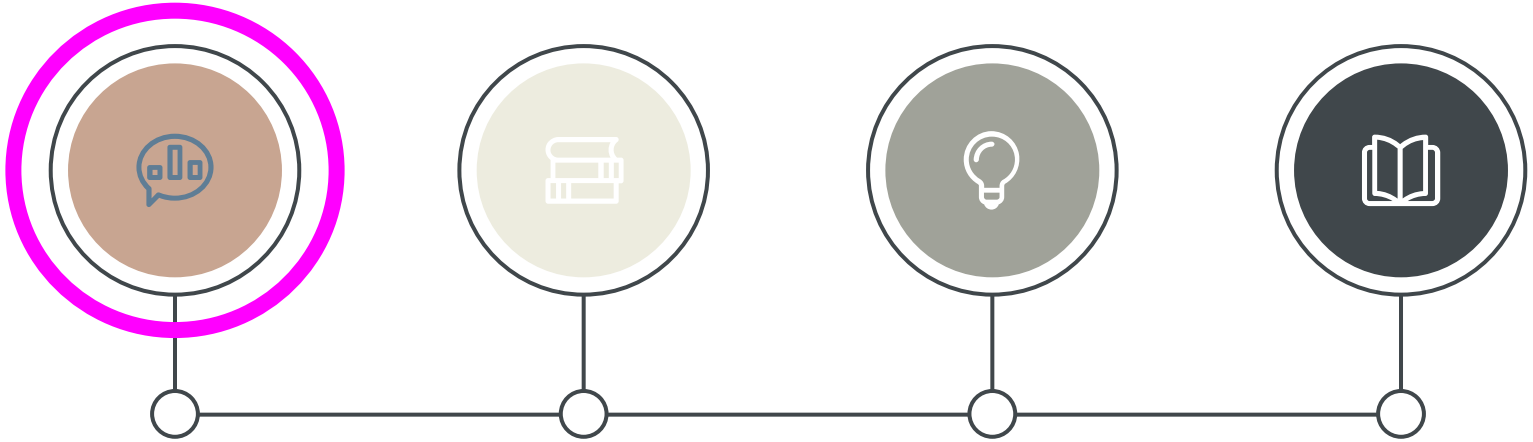

Data Cleaning

- Missing values
- Creating new features
- Choosing features

Modelling

- Building models
- Scoring models

Testing

- Kaggle testing on test dataset

Recommendations

- Important features

Missing Values

There are different types of missing values:

- Duplicate posts were removed
- Null and removed texts were replaced with the empty string
- One deleted post was miscategorized, we removed it

1969 rows left

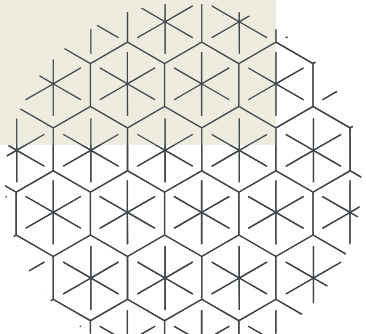

Plots of most frequent bigrams in the subreddits

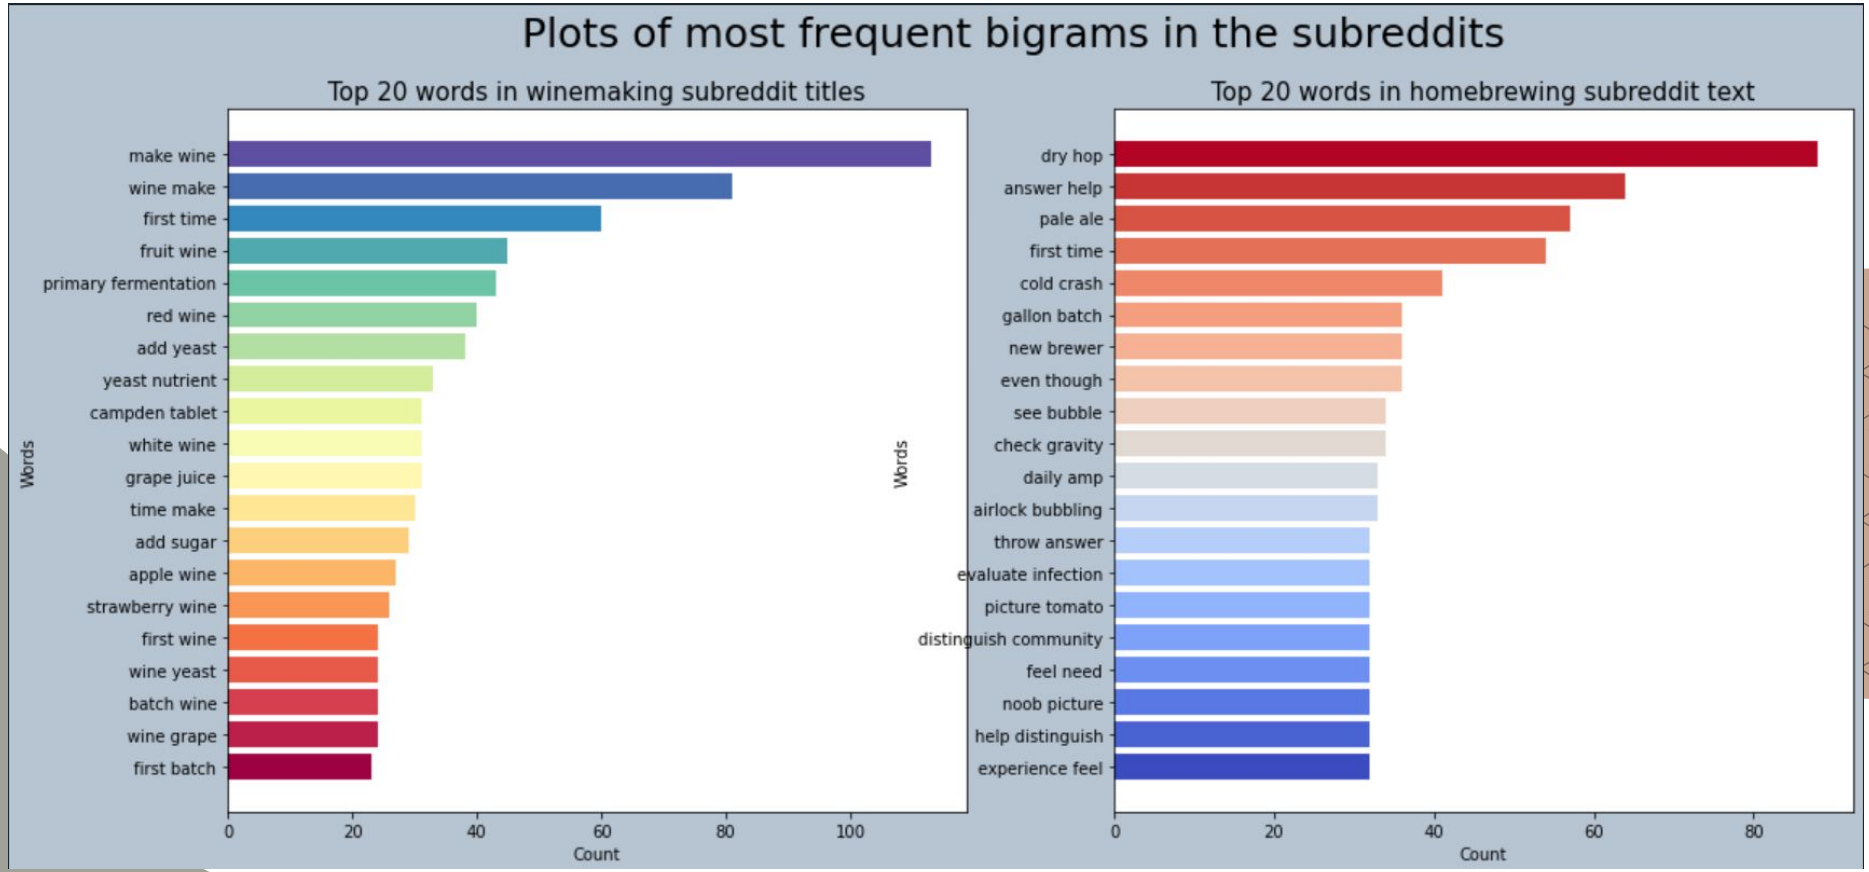

Plots of most frequent tf-idf bigrams in the subreddits

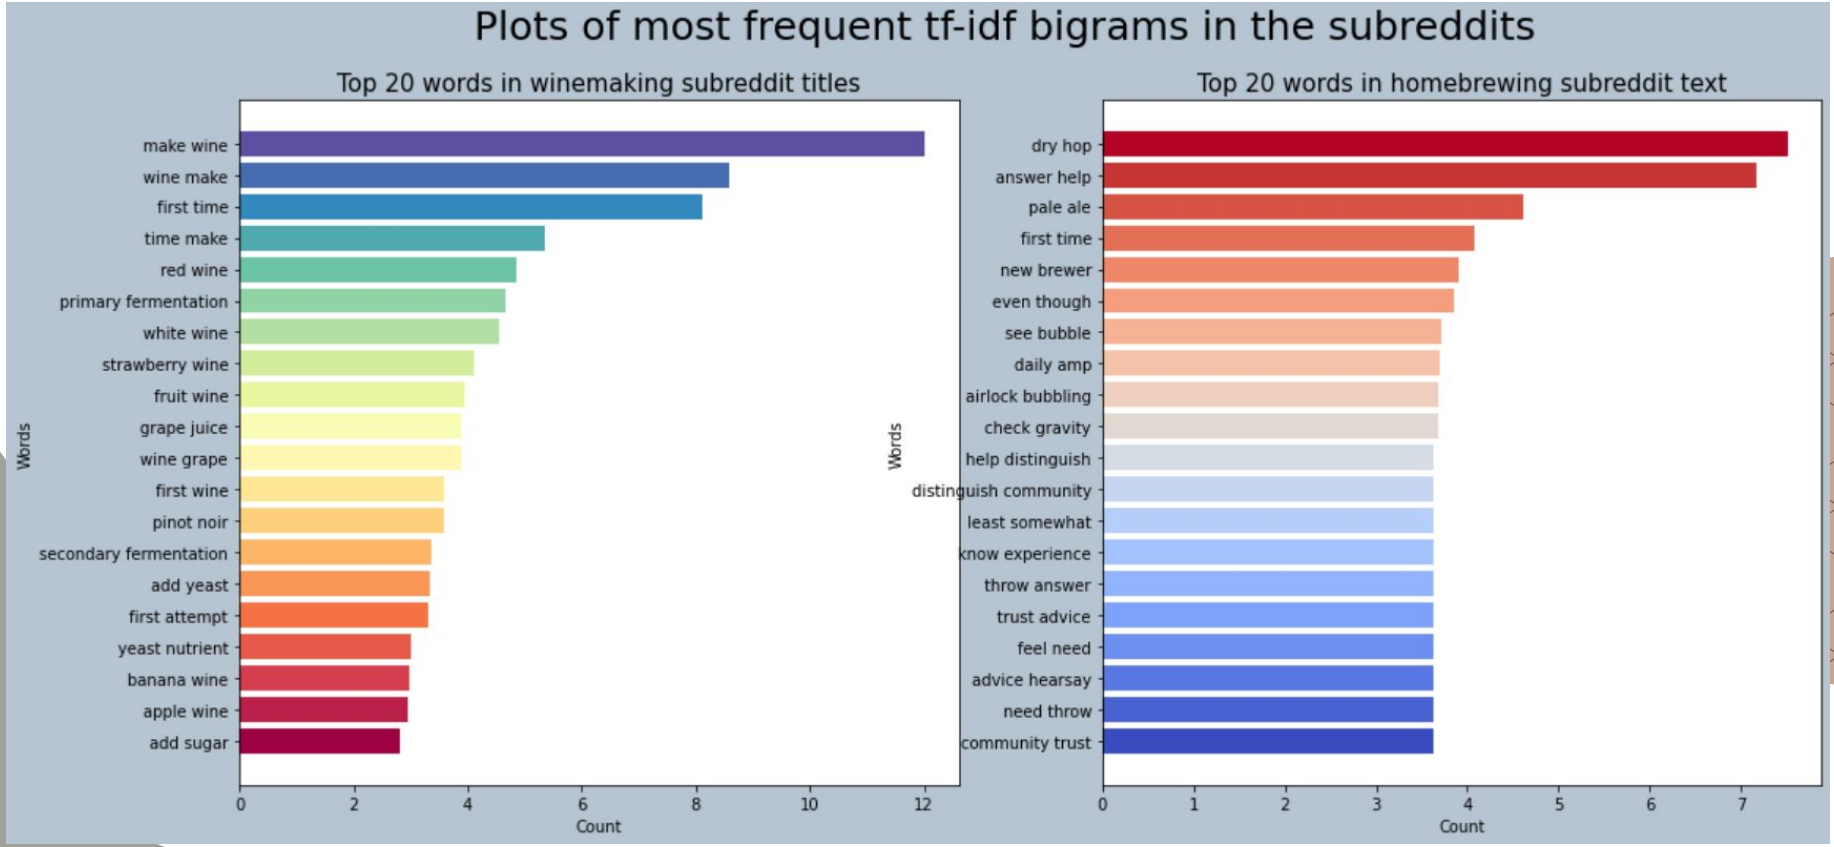

Model Workflow

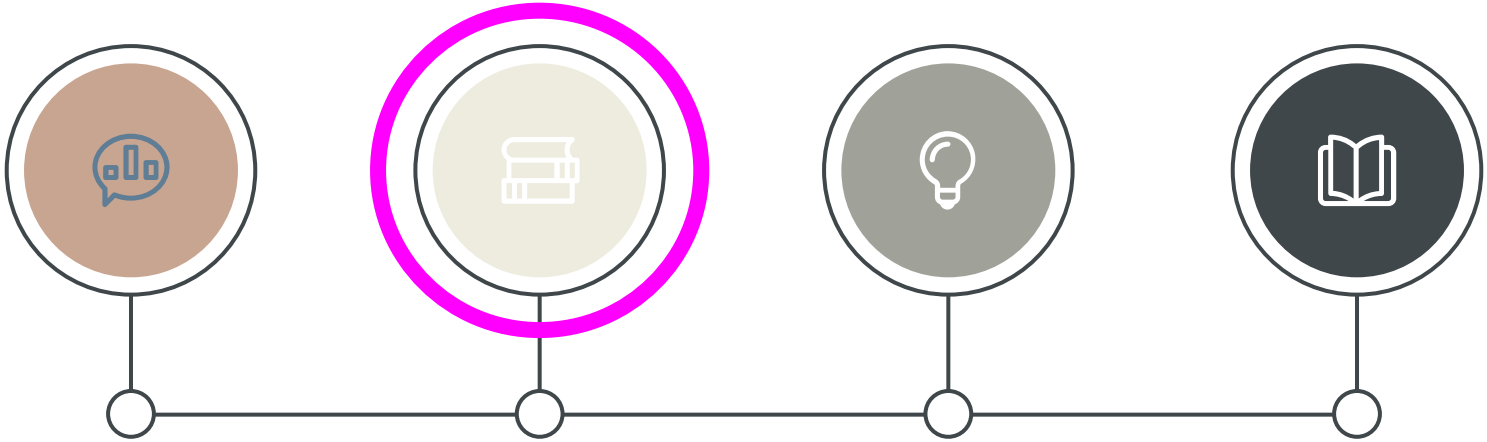

Data Cleaning

- Missing values
- Creating new features
- Choosing features

Modelling

- Building models
- Scoring models

Testing

- Kaggle testing on test dataset

Recommendations

- Important features

Modelling

- 8 models:
  - Logistic regression
    - Count
    - Tf-idf
  - KNN classifier
    - Count
    - Tf-idf
  - Naïve bayes
    - Count
    - Tf-idf
  - Random forest
    - Count
    - Tf-idf

Scoring

- ROC-AUC to determine best models
- F1 score to compare baseline score
- Accuracy to determine whether overfit

Model Workflow

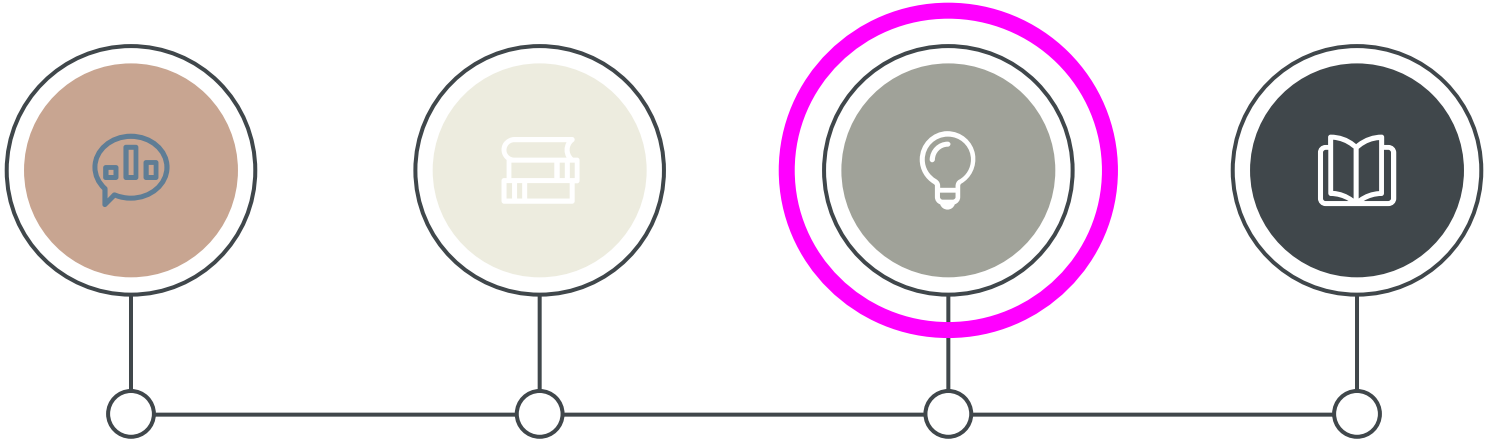

Data Cleaning

- Missing values
- Creating new features
- Choosing features

Modelling

- Building models
- Scoring models

Testing

- Kaggle testing on test dataset

Recommendations

- Important features

Logistic Regression performed the best

0.97 ROC-AUC
0.92 f1 score

models	vectorizer	f1 score	auc score
Logistic Regression	count	0.893	0.961
Logistic Regression	tf-idf	0.918	0.975
KNN Classifier	count	0.625	0.851
KNN Classifier	tf-idf	0.398	0.655
Naïve Bayes	count	0.749	0.925
Naïve Bayes	tf-idf	0.749	0.925
Random Forest	count	0.884	0.965
Random Forest	tf-idf	0.882	0.964

Model Workflow

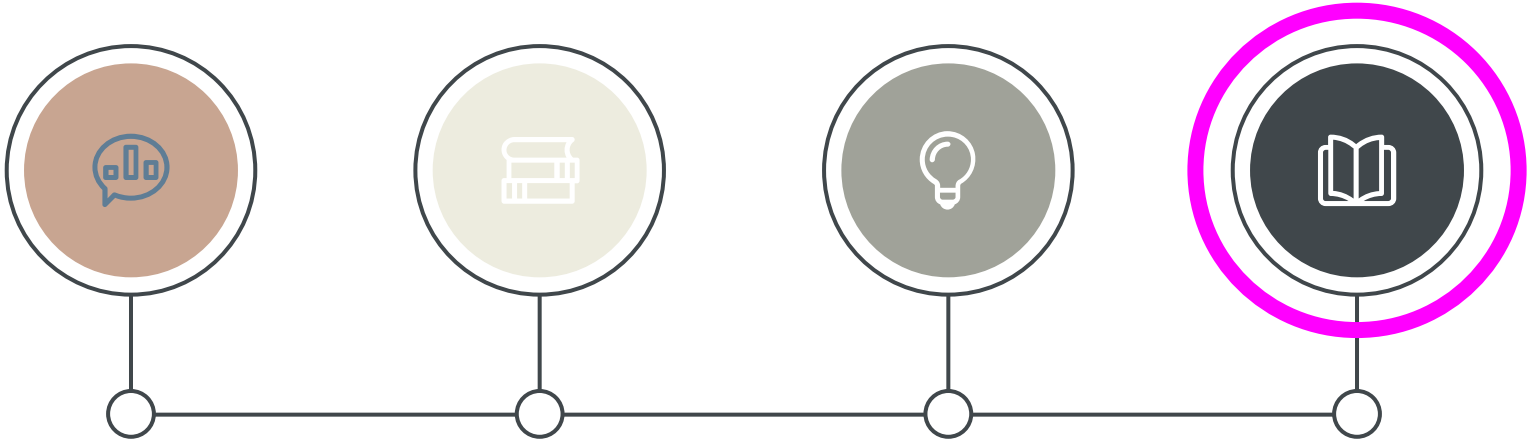

Data Cleaning

- Missing values
- Creating new features
- Choosing features

Modelling

- Building models
- Scoring models

Testing

- Kaggle testing on test dataset

Recommendations

- Important features

Best Model

Features

- Wordnet lemmatizer
- Tf-idf vectorizer
- Logistic regression with ridge penalty

Limitations

- Spell check before lemmatizing
- Slightly overfit, could remove more stopwords

Thanks